Effat College of Engineering · Computer Science

CS3081 Artificial Intelligence
Assignment: Uncertainty and Bayesian Networks

Instructor: Dr. Naila Marir **Total:** 50 marks (5 exercises × 10 marks)

Name: Muhammad Jeewani Student ID: S24109261

Instructions

1. Show all your steps. If you write only the final number, you get at most half of the marks.
2. Write probabilities with 4 decimal places (for example 0.5367), unless the question says something else.
3. Exercises 1 to 4, and Exercise 5 parts (a) and (b), are solved by hand.
4. Exercise 5 part (c) is solved in Python with pgmpy in VS Code, like in the lab.
5. Submit two files on Blackboard: these pages with your answers written in the boxes, scanned as one PDF, and your notebook `assignment_bn.ipynb` with all outputs.
6. This is individual work. You can discuss ideas with classmates, but you must write your own answers.

Exercise 1 — Joint distribution (10 marks)

The Effat shuttle bus is either **on time** or **late**. The weather in Jeddah is either **clear** or **dusty**. The table gives the probability of each pair. For example, the probability that the weather is clear **and** the bus is late is 0.14.

Weather	Shuttle = on time	Shuttle = late
clear	0.56	0.14
dusty	0.12	0.18

1. **(2 marks)** Show that the four numbers in the table make a valid probability distribution.

(0.56+0.14) + (0.12+0.18) = 0.7+0.3=1

2. **(2 marks)** Find P(dusty), the probability that the weather is dusty. Then find P(late), the probability that the bus is late.

P(dusty)=P(dusty AND on time)+ P(dusty AND late)= 0.12 + 0.18 = 0.3

P(late)=P(clear AND late)+ P(dusty AND late)= 0.14 + 0.18 = 0.32

3. **(2 marks)** Find P(late | dusty) and P(late | clear).

P(late | dusty)= (P(late AND dusty))/P(dusty)=0.18/0.3=0.6

4. **(2 marks)** Find P(dusty | late). Is it the same as P(late | dusty)? Explain your answer in one sentence.

P(dusty | late)= (P(late AND dusty))/P(late)=0.18/0.32=0.5625

Answers are different because they have different denominators

5. **(2 marks)** Are the weather and the bus independent? Use numbers to explain your answer.

P(clear) × P(late) = 0.7 × 0.32 = 0.224 ≠ 0.14 So not independent

---

Exercise 2 — Bayes' rule: a spam filter (10 marks)

An email program wants to find spam emails. It knows these facts:

- 30% of all emails are spam. So P(spam) = 0.3 and P(normal) = 0.7.
- In spam emails, the word "free" appears 60% of the time: P(free | spam) = 0.6.
- In normal emails, the word "free" appears 5% of the time: P(free | normal) = 0.05.

1. **(3 marks)** Find P(free), the probability that any email contains the word "free". Write the name of the rule you use.

P(free)= P(spam) * P(free | spam) + P(normal) * P(free | normal) =0.3(0.6)+0.7(0.05) = 0.215

2. **(3 marks)** A new email contains the word "free". Use Bayes' rule to find P(spam | free).

P(spam | free)= (P(spam) * P(free | spam))/P(free)= (0.3 * 0.6)/0.215= 0.8372

3. **(3 marks)** The program also checks the word "winner": P(winner | spam) = 0.4 and P(winner | normal) = 0.02. Assume that, once we know if an email is spam or normal, the two words appear independently of each other. A new email contains **both** "free" and "winner". Find P(spam | free, winner).

P(spam)= 0.3 × 0.6 × 0.4 = 0.0720
P(normal)= 0.7 × 0.05 × 0.02 = 0.0007
P(free, winner) = 0.0720+0.0007=0.0727
P(spam | free, winner) = 0.0720 / 0.0727 = 0.9904.

4. **(1 mark)** Compare your answers to parts 2 and 3. Did the probability of spam go up or down when we added the word "winner"? Explain why in one sentence.

It went up because “winner” is much more common in spam (0.4) than in normal email (0.02), so it is extra evidence for spam.

Exercise 3 — Bayesian network: structure and joint probability (10 marks)

In Jeddah summers, a **heatwave** (H) can cause an **AC overload** (A) in a campus building. A heatwave and an AC overload can both cause a **power cut** (C). A power cut can cause the **lab session to be cancelled** (L). Each variable is yes or no.

The network has 4 arrows: H → A, H → C, A → C, and C → L. The tables give only the probability of "yes". The probability of "no" is 1 minus the probability of "yes".

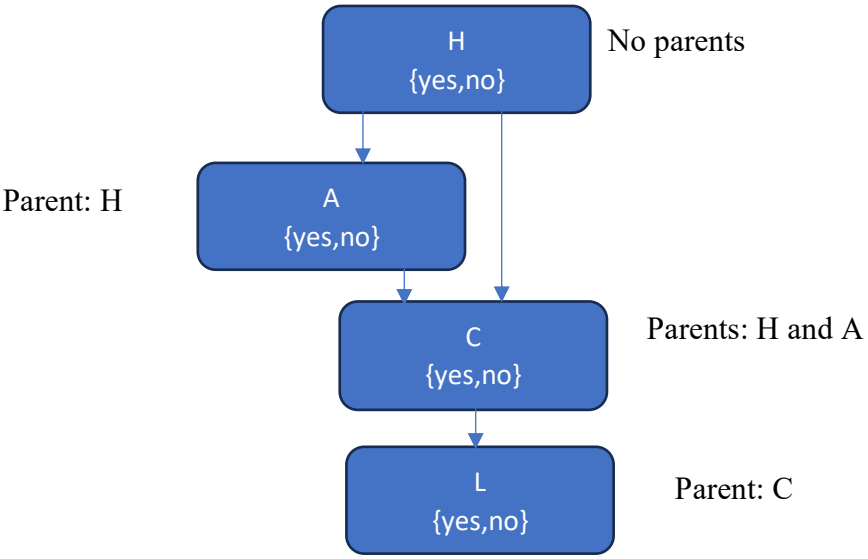
P(H = yes)
0.3

H	P(A = yes | H)
yes	0.6
no	0.1

H	A	P(C = yes | H, A)
yes	yes	0.7
yes	no	0.2
no	yes	0.5
no	no	0.05

C	P(L = yes | C)
yes	0.8
no	0.1

1. **(2 marks)** Draw the network. For each node, write its parents.



2. **(2 marks)** Write P(H, A, C, L) as a product of the probabilities from the network's tables.

P(H, A, C, L) = P(H) * P(A | H) * P(C | H, A) * P(L | C)

3. **(2 marks)** A full joint table for 4 yes/no variables needs 15 numbers ($2^4 - 1$). How many numbers do the 4 tables of this network need? Count only the "yes" probabilities, and show how you counted.

1 + 2 + 4 + 2 = 9 numbers

4. **(2 marks)** Find P(H = yes, A = no, C = yes, L = yes).

0.3 * 0.4 * 0.2 * 0.8 = 0.0192

5. **(2 marks)** Find P(H = yes | C = yes). Then give P(H = yes | C = yes, L = yes) without doing a new calculation, and explain why the two answers are the same.

P(H,C)=0.3*(0.6(0.7)+0.4(0.2))=0.15

P(-H, C) = 0.7(0.1·0.5 + 0.9·0.05) = 0.0665

P(C)=0.15+0.0665=0.2165 P(H|C)=0.15/0.2165=0.6928

P(H|C,L) will also be 0.6928 as L and H are independent

Exercise 4 — Inference by enumeration, challenge question (10 marks)

This is the **challenge question** of the assignment. Show every step.

Use the network and tables from Exercise 3. Today the lab session **was cancelled** (L = yes).

1. **(4 marks)** Use inference by enumeration to find P(H = yes | L = yes), the probability that there was a heatwave. Write the query variable, the evidence and the hidden variables, then show all sums and the normalization.

H	A	C	P(H)*P(A|H)*P(C|H,A)*P(L=yes|C)
yes	yes	yes	0.3*0.6*0.7*0.8 = 0.1008
yes	yes	no	0.3*0.6*0.3*0.1 = 0.0054
yes	no	yes	0.3*0.4*0.2*0.8 = 0.0192
yes	no	no	0.3*0.4*0.8*0.1 = 0.0096
no	yes	yes	0.7*0.1*0.5*0.8 = 0.0280
no	yes	no	0.7*0.1*0.5*0.1 = 0.0035
no	no	yes	0.7*0.9*0.05*0.8 = 0.0252
no	no	no	0.7*0.9*0.95*0.1 = 0.05985

Sum for H=yes: 0.135

Sum for H=no: 0.11655

Normalization: 0.135+0.11655=0.25155

P(H|L)=0.135/0.25155=0.5367

2. **(3 marks)** Now find P(A = yes | L = yes), the probability that there was an AC overload. Which variables are hidden this time?

Query: A. Evidence: L = yes. Hidden: H and C
A=yes terms: 0.1008 + 0.0054 + 0.0280 + 0.0035 = 0.1377
P(A=yes | L=yes) = 0.1377 / 0.25155 = 0.5474

3. **(3 marks)** A technician checks the building and reports that there was **no** AC overload (A = no). The lab is still cancelled (L = yes). Find P(H = yes | L = yes, A = no). Compare it with your answer to part 1: did the probability of a heatwave go up or down? Use the network to explain why in two or three sentences.

Now A=no is also evidence. H=yes: 0.0192 + 0.0096 = 0.0288. H=no: 0.0252 + 0.05985 = 0.08505. Total = 0.11385.
P(H=yes | L=yes, A=no) = 0.0288 / 0.11385 = 0.2530.
It went down, A heatwave makes an AC overload more likely, so learning there was no overload is evidence against a heatwave. It also removes one of the two routes by which a heatwave could have caused the power cut and the cancellation

Exercise 5 — Approximate inference by sampling (10 marks)

Now we find the same probability, P(H = yes | L = yes), by **sampling**. Sampling gives an estimate, not the exact answer.

(a) Rejection sampling (3 marks). Reminder: in rejection sampling, we keep only the samples that agree with the evidence (here L = yes), and we throw away the others. Below are 12 samples from the network of Exercise 3.

#	H	A	C	L
1	no	no	no	no
2	yes	yes	yes	yes
3	no	no	no	no
4	no	yes	yes	yes
5	yes	yes	no	no
6	no	no	no	yes
7	yes	no	yes	yes
8	no	no	no	no
9	no	no	no	no
10	yes	yes	yes	no
11	no	no	no	no
12	no	no	yes	yes

1. Write the numbers of the samples we keep and the samples we throw away. Then estimate P(H = yes | L = yes).

Keep: 2, 4, 6, 7, 12

Throw away: 1, 3, 5, 8, 9, 10, 11

H = yes among kept: samples 2 and 7, so the estimate is 2/5 = 0.4000

2. What fraction of the 12 samples did we throw away? Why is this a problem for rejection sampling?

7/12 = 0.5833 were thrown away. This is wasteful, and it gets much worse when the evidence is rare, because most samples are discarded and few remain for the estimate

(b) Likelihood weighting (3 marks). Reminder: in likelihood weighting, we do not sample the evidence. We fix it (here L = yes), we sample the other variables, and we give each sample a **weight**. The weight is the probability of the evidence given the sample's values: here, P(L = yes | C). Below are 6 samples.

#	H	A	C	L
1	no	no	no	yes (fixed)
2	yes	yes	yes	yes (fixed)
3	no	no	no	yes (fixed)
4	yes	no	no	yes (fixed)
5	no	yes	yes	yes (fixed)
6	yes	yes	yes	yes (fixed)

1. Write the weight of each sample.

Weight = P(L=yes | C), which is 0.8 if C = yes and 0.1 if C = no

Sample	**1**	**2**	**3**	**4**	**5**	**6**
Weight	0.1	0.8	0.1	0.1	0.8	0.8

2. Estimate P(H = yes | L = yes) using the weights: add the weights of the samples where H = yes, then divide by the sum of all weights.

H = yes samples are 2, 4, 6: 0.8 + 0.1 + 0.8 = 1.7. Total weight = 2.7. Estimate = 1.7 / 2.7 = 0.6296

3. Compare your estimates from (a) and (b) with the exact answer from Exercise 4. Why are they different? How can we make them closer to the exact answer?

Exact = 0.5367. The estimates (0.4000 and 0.6296) differ because they come from very few random samples, so they have high variance. Using many more samples brings them closer. Likelihood weighting also wastes no samples, so it usually converges faster.

(c) Python with pgmpy (4 marks). Answer this part in your notebook, not on paper. In VS Code, create `assignment_bn.ipynb`.

1. Build the network of Exercise 3 with `DiscreteBayesianNetwork` and `TabularCPD`, like in the lab. Name the nodes `heatwave`, `overload`, `powercut` and `cancel`, and use the values `yes` and `no`. Print `model.check_model()`. It must print `True`.
2. Use `VariableElimination` to find P(H | L = yes). Your answer must be the same as in Exercise 4.
3. Use rejection sampling with 10,000 samples to estimate P(H = yes | L = yes). Hint: import `BayesianModelSampling` from `pgmpy.sampling` and `State` from `pgmpy.factors.discrete`. Then use `BayesianModelSampling(model).rejection_sample(evidence=[State("cancel", "yes")], size=10000)`. This gives a table of samples. Count the fraction of rows where `heatwave` is `yes`.
4. In a Markdown cell, write two sentences that compare the exact answer and the sampling estimate.

